

AI Model Provider Comparison

Models released in the last 12 months (July 2025 – July 2026)

Model Provider	Model	Release Date	Cutoff Date	Cost of Tokens (per 1M, Input/Output)	Context Window	Type	When to Use This Model
OpenAI	GPT-5.5	Apr 24, 2026	~Apr 2026	\$5.00 / \$30.00 (Pro: \$30/\$180)	1,050,000 tokens (922K in / 128K out)	Reasoning, multimodal (text+image)	Complex professional workloads: agentic coding, research, data analysis, computer use, multi-step tasks with ambiguity
Google	Gemini 3.1 Pro	Feb 19, 2026	~early 2026	\$2.00 / \$12.00 (≤200K); \$4.00/\$18.00 (>200K)	1M input / 64K output tokens	Multimodal reasoning (text, image, audio, video, code)	Agentic workflows, algorithm design, large-scale data synthesis, sophisticated coding, full-repo code understanding
Anthropic	Claude Sonnet 5	Jun 30, 2026	Jan 2026	Intro: \$2/\$10 (through Aug 31, 2026); Standard: \$3/\$15	1M-token context, 128K max output (up to 300K via batch beta)	Agentic, hybrid reasoning	Customer-facing agents, production coding, computer use, high-volume professional workflows (finance, research, content)
xAI	Grok 4.3	Apr 30, 2026	Dec 2025	\$1.25 / \$2.50	1M-token context window	Multi-agent reasoning, video input	Native video input, document generation (PDFs/spreadsheets/slides), improved tool-calling, real-time X/web grounding
Meta	Muse Spark (Meta Superintelligence Labs)	Apr 8, 2026	~2026	Not publicly listed (API-only)	Not specified in available sources	Closed-weight reasoning model	First proprietary frontier model from Meta; leads HealthBench Hard vs Gemini 3.1 Pro – health/reasoning-heavy tasks
DeepSeek	DeepSeek V4 (Pro/Flash)	Apr 24, 2026 (preview)	~2026	Not fully public yet	1M-token native context; V4-Pro: 1.6T total/49B active params; V4-Flash: 284B total/13B active	Open-weight MoE, hybrid attention	Long-context tasks at much lower inference cost (27% FLOPs, 10% KV cache vs V3.2 at 1M context)
Mistral AI	Mistral Large 3	2025 (exact date unconfirmed)	~2025	Not confirmed in sources	Not confirmed (predecessor Large 2.1 had 128K)	Dense open-weight (Apache 2.0)	European enterprises needing GDPR-compliant, multilingual solutions
Alibaba	Qwen3.7 Max	May 19–20, 2026	~2026	\$2.50 / \$7.50 (90% cached-input discount)	1M-token context window	Closed, agent-first flagship	Long-horizon coding agents, hours-long tasks with hundreds/thousands of tool calls, office automation, deep research agents
Amazon	Nova 2 (family)	Announced Dec 2025 (AWS re:Invent)	~2025	Via Amazon Bedrock (varies by tier)	Not detailed in available sources	Multi-tier foundation models, Bedrock-exclusive	Enterprise workloads needing AWS-native deployment across price/performance tiers
Cohere	Command A Reasoning	Aug 2025	~2025	~\$1.56 / \$1.56 (varies by host)	256,000-token context, up to 32,000 output tokens	Hybrid reasoning, agentic	Comprehensive Q&A, intricate document review, structured arguments, ReAct agents, RAG with citations